

A Deep Dive Into eBPF Program Loader

Cong Wang xiyou.wangcong@gmail.com

Linux Kernel Maintainer

Open Source Summit North America, 2025

Agenda

1. **eBPF Objects & Syscall Interface**
2. **Four-Phase Loading Pipeline**
3. **BTF Type System & CO-RE**
4. **Kernel Verifier Interaction**
5. **Program Signing Challenges**

eBPF Code Example

```
// Map definition (demo)
struct {
    __uint(type, BPF_MAP_TYPE_HASH);
    __uint(max_entries, 1024);
    __type(key, int);
    __type(value, struct event);
} events SEC(".maps");

// Program definition
SEC("kprobe/sys_open")
int trace_open(struct pt_regs *ctx) {
    bpf_printk("Hello world!");
    return 0;
}
```

eBPF Objects: Programs and Maps

1. Programs

- Executable bytecode
- Verified by kernel
- Attached to kernel hooks
- Type-specific (XDP, kprobe, etc.)

2. Maps

- Data structures for storage
- Shared between programs/userspace
- Various types (hash, array, etc.)
- Persistent across program runs

eBPF Objects: BTF and Links

3. BTF (BPF Type Format)

- Type information metadata
- Enables CO-RE and debugging
- Describes program/map structures
- Kernel and userspace type matching

4. Links

- Connection between program & hook
- Manages attachment lifecycle
- Automatic cleanup on process exit
- Reference counting for sharing

eBPF Object Creation via Syscall

```
// Create a map
map_fd = bpf(BPF_MAP_CREATE, &map_attr, sizeof(map_attr));
// Load a program
prog_fd = bpf(BPF_PROG_LOAD, &prog_attr, sizeof(prog_attr));
// Load BTF type information
btf_fd = bpf(BPF_BTF_LOAD, &btf_attr, sizeof(btf_attr));
// Create attachment link
link_fd = bpf(BPF_LINK_CREATE, &link_attr, sizeof(link_attr));
```

Key Insight: File descriptors (fd) are the kernel's handle to loaded objects

File Descriptor Management

- Each object gets a unique file descriptor
- FDs enable object sharing and persistence
- Objects destroyed when all FDs closed
- Can be pinned to filesystem for persistence
- **Challenge:** FDs are unpredictable at compile time

From Source Code to ELF Binary

Compilation Process:

```
# Compile eBPF C source to ELF object file  
clang -target bpf -O2 -c program.c -o program.o
```

What happens during compilation:

- eBPF C source code → eBPF bytecode instructions
- Metadata (maps, BTF) → ELF sections
- Result: ELF binary containing **only** eBPF programs and metadata
- libbpf is **not** linked into the eBPF binary
- libbpf is a userspace library linked into **your** loader

ELF Code and Data Sections

- `SEC("type/name")` - Main program bytecode
- `.text` - Subprogram/helper function instructions
- `.rodata` - Read-only data (strings, constants)
- `.data` - Initialized global variables
- `.bss` - Uninitialized global variables

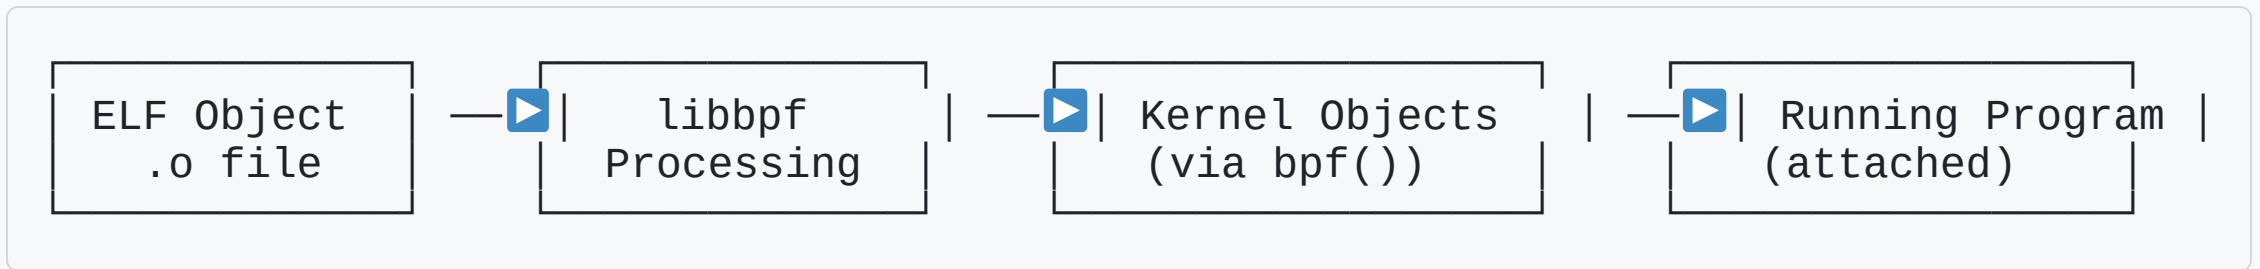
ELF Metadata Sections

- `.maps` - Map definitions and attributes
- `.BTF` - Type information (structs, functions)
- `.BTF.ext` - Source line info, function info
- `.rel*` - Relocation information

ELF to Kernel Object

- Step 1: ELF Parsing
 - `obj = bpf_object__open("program.o");`
 - Parse ELF file, extracts sections, symbols, relocations
- Step 2: Object Preparation
 - `bpf_object__prepare(obj);`
 - Resolves symbols, creates maps, processes relocations
- Step 3: Program Loading
 - `bpf_object__load(obj);`
 - Each program loaded and verified for safety and correctness
- Step 4: Attachment and Execution

The Complete Loading Process



Key Insight: libbpf acts as a sophisticated translator between static ELF representation and dynamic kernel objects

Core Data Structures: `struct bpf_object`

```
enum bpf_object_state {
    OBJ_OPEN,          // ELF parsed, sections identified
    OBJ_PREPARED,       // Maps created, relocations done
    OBJ_LOADED          // Programs loaded to kernel
};

struct bpf_object {
    char name[BPF_OBJ_NAME_LEN];          // Object name
    char license[64];                     // License string
    enum bpf_object_state state;           // Loading state

    struct bpf_program *programs;         // Program array
    struct bpf_map *maps;                 // Map array
    struct btf *btf;                      // Object BTF
    struct btf *btf_vmlinux;              // Kernel BTF

    int *fd_array;                        // File descriptor storage
    size_t fd_array_cnt;                  // Number of FDs
};
```

Core Data Structures: Programs & Maps

Program Representation

```
struct bpf_program {  
    char *name;                // Program name  
    enum bpf_prog_type type;    // XDP, kprobe, etc  
    struct bpf_insn *insns;     // Instructions  
    size_t insns_cnt;          // Instruction count  
  
    int fd;                    // Kernel FD after load  
    struct reloc_desc *reloc_desc; // Relocations  
    int nr_reloc;              // Relocation count  
};
```

Map Definition

```
struct bpf_map {  
    char *name;                // Map name  
    enum bpf_map_type type;     // HASH, ARRAY, etc  
    __u32 key_size, value_size; // Entry sizes  
    __u32 max_entries;          // Capacity  
  
    int fd;                    // Kernel FD after create  
    __u32 btf_key_type_id;      // BTF type IDs  
    __u32 btf_value_type_id;  
};
```

Phase 1 - Discovery

ELF Section Processing:

```
bpf_object__elf_collect() // Parse ELF sections (.text, .maps, .BTF)
bpf_object__add_programs() // Extract program instructions from .text
bpf_object__init_user_btf_maps() // Parse map definitions from .maps
```

BTF and Metadata:

```
bpf_object__init_btf() // Load BTF type information
```

Output: Parsed object with identified sections, programs, and maps

Phase 2 - Resolution

Kernel Interaction Setup:

```
bpf_object_prepare_token() // Prepare for kernel interaction  
btf__load_vmlinux_btf() // Load kernel BTF for type matching
```

Symbol Resolution & Relocations:

```
bpf_object__resolve_externs() // Resolve external symbols (kernel config, functions)  
bpf_object__relocate() // Process all relocations
```

Map Creation

```
bpf_object__create_maps() // Create eBPF maps with BPF_MAP_CREATE
```

Output: Fully resolved objects ready for kernel loading

Phase 3 - Kernel Interaction

Program Loading Loop:

```
bpff_object__load_progs(obj, extra_log_level) {  
    for (i = 0; i < obj->nr_programs; i++) {  
        bpff_object_load_prog(obj, prog, prog->insns, prog->insns_cnt,  
                               obj->license, obj->kern_version, &prog->fd);  
    }  
}
```

Verifier Interaction:

```
fixup_verifier_log() // Post-process verifier log to improve error descriptions
```

Output: Loaded programs with kernel file descriptors

Phase 4 - Attachment

Program Attachment Methods:

```
// Auto-attachment based on SEC() definition
bpf_object__attach_skeleton(skel);
// Manual attachment for specific hooks
link = bpf_program__attach_kprobe(prog, false, "sys_open");
link = bpf_program__attach_xdp(prog, ifindex);
link = bpf_program__attach_cgroup(prog, cgroup_fd);
```

Execution Context: Program runs when kernel events trigger the attached hook

Link Management

- Links represent the attachment relationship
- Can be pinned to filesystem for persistence
- Automatically cleaned up when program exits
- Support for multiple programs on same hook

BTF: The Type System Foundation

BTF Format

- Binary format encoding C type information
- Compact representation with string deduplication
- Enables CO-RE relocations and type verification

Kernel BTF Registry

- vmlinux BTF: Core kernel types
- Module BTF: Per-module types
- User BTF: Object-specific types
- Type ID remapping for compatibility

Map FD Relocation

- Compiler generates placeholders in map access instructions
- Libbpf parses `.maps` definition and creates maps
- Libbpf stores FDs in `obj->fd_array[]`
- Program instructions are patched with FD values
- The verifier will translate FD to map pointer

Relocation Engine

CO-RE Relocation Types:

```
enum bpf_core_relo_kind {  
    BPF_RELO_FIELD_BYTE_OFFSET,    // struct field access  
    BPF_RELO_FIELD_BYTE_SIZE,     // sizeof operations  
    BPF_RELO_TYPE_ID_TARGET,       // kernel type matching  
    BPF_RELO_ENUMVAL_VALUE,        // enum value resolution  
};
```

Internal Relocation Types:

```
enum reloc_type {  
    RELO_LD64,    // Map loading (FD → instruction)  
    RELO_CALL,    // Function calls  
    RELO_DATA,    // Global data access  
    RELO_EXTERN_LD64, // External symbols  
    RELO_CORE,    // CO-RE relocations  
};
```

Example: Map Relocation

Map Reference in Program:

```
// BPF program referencing a map
struct bpf_map_def SEC("maps") my_map = { /* ... */ };

SEC("kprobe/sys_open")
int trace_open(void *ctx) {
    bpf_map_lookup_elem(&my_map, &key); // This generates RELO_LD64 relocation
}
```

Relocation Processing:

```
case RELO_LD64: // Map FD relocation
    map_fd = obj->maps[relo->map_idx].fd;
    // Patch instruction with actual FD
    insn->src_reg = BPF_PSEUDO_MAP_FD;
    insn->imm = map_fd;
    break;
```

CO-RE (Compile Once, Run Everywhere)

The Problem:

- Kernel structures change across versions
- Field offsets differ between kernels
- Traditional eBPF programs break on kernel updates

CO-RE Solution

```
// In BPF program - symbolic access
struct task_struct *task = ...;
int pid = BPF_CORE_READ(task, pid); // Field offset unknown at compile time

// At load time - resolved via BTF by the kernel
case BPF_RELO_FIELD_BYTE_OFFSET:
    offset = btf_member_bit_offset(target_btf, member_idx);
    bpf_core_patch_insn(prog_name, insn, insn_idx, relo, relo_idx, &targ_res);
```

Benefits:

- Same binary runs on multiple kernel versions
- Type-safe field access with BTF validation
- Automatic adaptation to kernel structure changes

External Symbol Resolution

Kernel Feature Adaptation:

- Programs adapt behavior based on kernel capabilities
- Feature detection without runtime checks
- Compile-time optimization with runtime flexibility

Resolution Process

External Symbol Types:

- `EXT_KCFG` - Kernel config symbols (`CONFIG_*`)
- `EXT_KSYM` - Kernel symbols (`/proc/kallsyms`)
- `EXT_KFUNC` - Kernel function addresses
- `EXT_DATASEC` - Data section references

Resolution Process:

```
// Example: kernel config resolution
extern bool CONFIG_BPF_SYSCALL __kconfig; // in your eBPF code

// libbpf: read /proc/config.gz or /boot/config-*
bpf_object__read_kconfig_file(obj, kcfg_data); // Sets CONFIG_BPF_SYSCALL = 1 or 0
```

Kernel Verifier Interaction

BPF_PROG_LOAD Attributes:

```
union bpf_attr attr = {  
    .prog_type = prog->type,           // XDP, kprobe, etc  
    .insn_cnt = prog->insns_cnt,       // Instruction count  
    .insns = ptr_to_u64(prog->insns),  // Program bytecode  
    .log_size = log_buf_size,         // Log buffer size  
    .log_buf = ptr_to_u64(log_buf),    // Verifier output  
    .prog_btfd_fd = prog->btfd_fd,     // Type information  
};
```

Syscall and Response:

- `fd = syscall(__NR_bpf, BPF_PROG_LOAD, &attr);`
- Success: Returns program FD
- Failure: Returns -1, verifier log has details

Verifier Log Processing

BPF Verifier Visualizer (bpfvv):

- Interactive tool for debugging verification failures
- Parses verifier logs and visualizes program states
- Acts as primitive debugger UI for BPF programs
- Load text files containing verifier output for analysis
- Available at: <https://libbpf.github.io/bpfvv/>

Program Signing Challenges

Current State:

- No built-in signing mechanism in current libbpf and kernel

Key Challenges:

- Post-relocation signing complexity
- Kernel version compatibility
- Dynamic symbol resolution

Proposed Solutions:

- 2-phase Signing: <https://github.com/congwang/ebpf-2-phase-signing>
- Hornet LSM: light-skeleton-based eBPF signature verification

Key Takeaways

1. **Layered Design:** libbpf provides sophisticated abstraction over kernel complexity
2. **Type Safety:** BTF enables portable, type-safe program development
3. **Smart Relocation:** CO-RE technology enables true portability
4. **Robust Verification:** Multi-stage validation ensures program safety

Bottom Line: eBPF program loading is a carefully orchestrated dance between userspace tooling and kernel verification

Thank You

Key Resources:

- libbpf source: `github.com/libbpf/libbpf`
- eBPF documentation: `ebpf.io`
- Kernel BPF documentation: `kernel.org/doc/html/latest/bpf/`
- Kernel mailing list: `bpf@vger.kernel.org`